

Lab 2: Deep Q-Network (DQN)

Lab Objective:

In this lab, you will learn and implement the Deep Q-Network algorithm by solving MsPacman-v5.

Important Date:

- Submission deadline: 10/15 (Sun) 23:59

Turn in:

1. Experiment report (.pdf)
2. Source code [NOT including model weights]

Notice: zip all files with name “**RL_LAB2_StudentId_Name.zip**”,

e.g.: 「RL_LAB2_311551059_陳昱丞.zip」

RL_LAB2_311551059_陳昱丞

```
|— src
    |— base_agent.py
```

```
...
|— report.pdf
```

(Put all your source code in **src** directory)

(Wrong format deduction: -5pts; Multiple deductions may apply.)

Lab Description:

- Understand the mechanism of both behavior network and target network.
- Understand the mechanism of experience replay buffer.
- Learn to construct and design neural networks.

Requirements:

- Implement DQN
 - Construct the neural network.
 - Select action according to epsilon-greedy.
 - Construct Q-values and target Q-values.
 - Calculate loss function.
 - Update behavior and target network.
 - Understand Deep Q-learning mechanisms.
 - Understand common techniques in Atari environment like frame stack, grayscale...

Game Environment – MsPacman-v5:

- Introduction: Your goal is to collect all of the pellets on the screen while avoiding the ghosts.
- Observation: By default, the environment returns the RGB image that is displayed to human players as an observation.

- Action [9]:

Num	Action
0	NOOP
1	UP
2	RIGHT
3	LEFT
4	DOWN
5	UPRIGHT
6	UPLEFT
7	DOWNRIGHT
8	DOWNLEFT

Algorithm – Deep Q-learning with experience replay:

```

Initialize replay memory  $D$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights  $\theta$ 
Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
For episode = 1,  $M$  do
  Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$ 
  For  $t = 1, T$  do
    With probability  $\varepsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$ 
    Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the
    network parameters  $\theta$ 
    Every  $C$  steps reset  $\hat{Q} = Q$ 
  End For
End For

```

Implementation Details:

Network Architecture

- Deepmind DQN network. Please refer to the sample code for details.

Training Hyper-Parameters

- Memory capacity (experience buffer size): 100000
- Batch size: 32
- Warmup steps: 20000
- Optimizer: Adam
- Learning rate: 0.0000625
- Epsilon: $1 \rightarrow 0.1$
- Gamma (discount factor): 0.99
- Update network every 4 steps
- Update target network every 10000 steps

You can tune the hyperparameter yourself.

Bonus:

- Solve Enduro-v5 using DQN.
- Add three improvements in your DQN implementation.
 1. Double DQN.
 2. Dueling DQN.
 3. DQN with parallelized rollout.

Scoring Criteria:

Show your results, otherwise no credit will be granted.

Your Score = report (30%) + report bonus (20%) + demo performance (50%) + demo questions (20%)


- **Report contains two parts:**
 - **Experimental Results (30%)**
 - (1) Screenshot of Tensorboard training curve and testing results on DQN.
 - **Experimental Results and Discussion of bonus parts (bonus) (20%)**
 - (1) Screenshot of Tensorboard training curve and testing results on Enduro-v5 (10%).
 - (2) Screenshot of Tensorboard training curve and testing results on DDQN, and discuss the difference between DQN and DDQN (3%).
 - (3) Screenshot of Tensorboard training curve and testing results on Dueling DQN, and discuss the difference between DQN and Dueling DQN (3%).
 - (4) Screenshot of Tensorboard training curve and testing results on DQN with parallelized rollout, and discuss the difference between DQN and DQN with parallelized rollout (4%).
- **Demo Performance (50%):**
 - Test your best model for one game.
 - You have to show the video while testing. You can use `env.render()` or save video function to achieve this.
 - You can use a fixed random seed to reproduce your best game score.
 - **Demo performance – Score table:**



Reward (Game score)	Points (50%)
0 ~ 500	0
500 ~ 1000	10
1000 ~ 1500	20
1500 ~ 2000	30
2000 ~ 2500	40
2500 ~ 3000	45
> 3000	50

Examples of Tensorboard training curve and testing results:


- Training curve:



- Testing results (5 games):

```
Episode: 1      Length: 13301   Total reward: 1040.00
Episode: 2      Length: 13307   Total reward: 1066.00
Episode: 3      Length: 9969    Total reward: 705.00
Episode: 4      Length: 13286   Total reward: 1059.00
Episode: 5      Length: 16630   Total reward: 1318.00
average score: 1037.6
```

- Training curve (comparison):



References:

- [1] Mnih, Volodymyr et al. “Playing Atari with Deep Reinforcement Learning.” ArXiv abs/1312.5602 (2013).
- [2] Mnih, Volodymyr et al. “Human-level control through deep reinforcement learning.” Nature 518 (2015): 529-533.
- [3] Van Hasselt, Hado, Arthur Guez, and David Silver. “Deep Reinforcement Learning with Double Q-Learning.” AAAI. 2016.
- [4] Lillicrap, Timothy P. et al. “Continuous control with deep reinforcement learning.” CoRR abs/1509.02971 (2015).
- [5] Silver, David et al. “Deterministic Policy Gradient Algorithms.” ICML (2014).
- [6] OpenAI. “OpenAI Gym Documentation.” Retrieved from Getting Started with Gym: <https://gym.openai.com/docs/>.
- [7] PyTorch. “Reinforcement Learning (DQN) Tutorial.” Retrieved from PyTorch Tutorials: https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html.