

Introduction to Reinforcement Learning (RL)

I-Chen Wu

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998.
 - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
 - Bible in this area.
- David Silver, Online Course for Deep Reinforcement Learning.
 - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>



Successful RL Examples

- Games: Super-human levels
 - Backgammon (Tesauro, 1994).
 - Connect6/2048/Threes! (CGI, 2022). Reach the top levels.
 - AlphaGo/AlphaZero/Muzero, using deep reinforcement learning (2016)
 - Open AI Five for Dota 2, 2019
 - AlphaStar for StarCraft by DeepMind (in nature), 2019
- Robotics: robot-controlled helicopters and humanoid robot walk (Abbeel et al.).
- Autonomous driving/racing: AWS DeepRacer (Amazon, CGI, 2019-)
- Manufacturing scheduling (CGI, 2022).
- Chip design: a fast graph placement by Google Brain (Nature, 2021)
- Optimizing matrix multiplication: AlphaTensor (2022)
- Chat bot: RLHF in Chat-GPT (OpenAI, 2022)
 - Reinforcement Learning from Human Feedback
- ...(Many more successful examples for deep reinforcement learning)

Stochastic Game: 2048

2	32768	8192	4096
16384	1024	512	256
2048	32	64	128
16	16	2	4

The First Game Reaching 65536 in the World (in 10,000 Trials) in 2015

<http://2048.aigames.nctu.edu.tw/replay.php>

2048

SCORE 1031392 BEST 1031392

512	256	32	2
1024	128	16	4
4096	64	8	2
65536	4	2	4

Game over!

Try again

AlphaGo/AlphaZero

● The Game of Go

- AlphaGo vs. 李世石: 4:1 (2016)
- AlphaGo Zero vs. 柯潔: 3:0 (2017)



Dota 2: OpenAI Five (OpenAI)


- Dota 2 is played with two teams defending bases in opposite corners. Each team have five players, each controlling a hero unit with unique abilities
- OpenAI Five became the first AI system to defeat the world champions at an esports game (2019)



Source of image: <https://technews.tw/2019/04/16/ai-dota-fight-alongside/>
Berner, Christopher, et al. "Dota 2 with large scale deep reinforcement learning." *arXiv preprint arXiv:1912.06680* (2019).

StarCraft II: AlphaStar (DeepMind)

- StarCraft is a real-time strategy game in which players balance high-level economic decisions with individual control of hundreds of units
- AlphaStar** was rated at Grandmaster level for all three StarCraft races, above 99.8% of officially ranked human players (2019)



Pluribus Poker (CMU and FAIR)

- Pluribus: stronger than top human professionals in six-player no-limit Texas hold'em poker (2019)




Source of image: <https://zhuanlan.zhihu.com/p/73336511>

Noam Brown Tuomas Sandholm, Superhuman AI for multiplayer poker. Science 365, 885-890 (2019). DOI: 10.1126/science.aay2400



AlphaTensor (2022)

Size (n, m, p)	Best method known	Best rank known	AlphaTensor rank	
			Modular	Standard
(2, 2, 2)	(Strassen, 1969) ²	7	7	7
(3, 3, 3)	(Laderman, 1976) ¹⁵	23	23	23
(4, 4, 4)	(Strassen, 1969) ² (2, 2, 2) \otimes (2, 2, 2)	49	47	49
(5, 5, 5)	(3, 5, 5) + (2, 5, 5)	98	96	98
(2, 2, 3)	(2, 2, 2) + (2, 2, 1)	11	11	11
(2, 2, 4)	(2, 2, 2) + (2, 2, 2)	14	14	14
(2, 2, 5)	(2, 2, 2) + (2, 2, 3)	18	18	18
(2, 3, 3)	(Hopcroft and Kerr, 1971) ¹⁶	15	15	15
(2, 3, 4)	(Hopcroft and Kerr, 1971) ¹⁶	20	20	20
(2, 3, 5)	(Hopcroft and Kerr, 1971) ¹⁶	25	25	25
(2, 4, 4)	(Hopcroft and Kerr, 1971) ¹⁶	26	26	26
(2, 4, 5)	(Hopcroft and Kerr, 1971) ¹⁶	33	33	33
(2, 5, 5)	(Hopcroft and Kerr, 1971) ¹⁶	40	40	40
(3, 3, 4)	(Smirnov, 2013) ¹⁸	29	29	29
(3, 3, 5)	(Smirnov, 2013) ¹⁸	36	36	36
(3, 4, 4)	(Smirnov, 2013) ¹⁸	38	38	38
(3, 4, 5)	(Smirnov, 2013) ¹⁸	48	47	47
(3, 5, 5)	(Sedoglavic and Smirnov, 2021) ¹⁹	58	58	58
(4, 4, 5)	(4, 4, 2) + (4, 4, 3)	64	63	63
(4, 5, 5)	(2, 5, 5) \otimes (2, 1, 1)	80	76	76



Alhussein Fawzi, Matej Balog, Aja Huang, Thomas Hubert, Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Francisco J. R. Ruiz, Julian Schrittwieser, Grzegorz Swirszcz, David Silver, Demis Hassabis, and Pushmeet Kohli. Discovering faster matrix multiplication algorithms with reinforcement learning. *Nature*, 610:47–53, 2022.



DeepRacer


- AWS DeepRacer (by our lab CGI)
 - 2020 AWS DeepRacer World Championship Cup: 1st + 3rd places
 - 2022 AWS DeepRacer World Championship Cup: 1st + 2nd + 3rd places



Better and Faster Chip Design

- Better and faster for chip design than any human designer.
 - Generate chip floorplans that are comparable or superior to human experts in under six hours,
 - whereas humans take months to produce acceptable floorplans for modern accelerators.

[1] A. Mirhoseini, et al. (by Google brain), A graph placement methodology for fast chip design, Nature, 2021




Reinforcement Learning from Human Feedback (RLHF) for ChatGPT

By OpenAI CEO (2022)

(at 6:56/2:23:56, Sam Altman in Lex Fridman Podcast)

- “... And RLHF is how we take some human feedback,
 - the simplest version of this is show two outputs
 - ask which one is better than the other
 - which one the human raters prefer
 - and then feed that back into the model with RL
 - **that process works remarkably well with in my opinion**
 - **remarkably little data to make the model more useful**
- So, RLHF is how we align the model to what humans want it to do.
...”



David Silver:
(the leader of the AlphaGo team)

“DL+RL = AI”



Many Faces of Reinforcement Learning

- Computer Science
 - Machine Learning
- Engineering
 - Optimal Control
- Mathematics
 - Operations Research
- Economics
 - Bounded Rationality
- Psychology
 - Classical/Operant Conditioning
- Neuroscience
 - Reward System



What are different from others?

- Characteristics:

- No supervisor, only a **reward** signal
- Feedback is delayed, not instantaneous
- Time really matters
- Agent's actions affect the subsequent data and actions

- UL vs. RL:

- RL is learning from interaction.
- RL does not rely on examples of correct behavior.
- RL is trying to maximize a reward signal, instead of trying to find hidden structure.



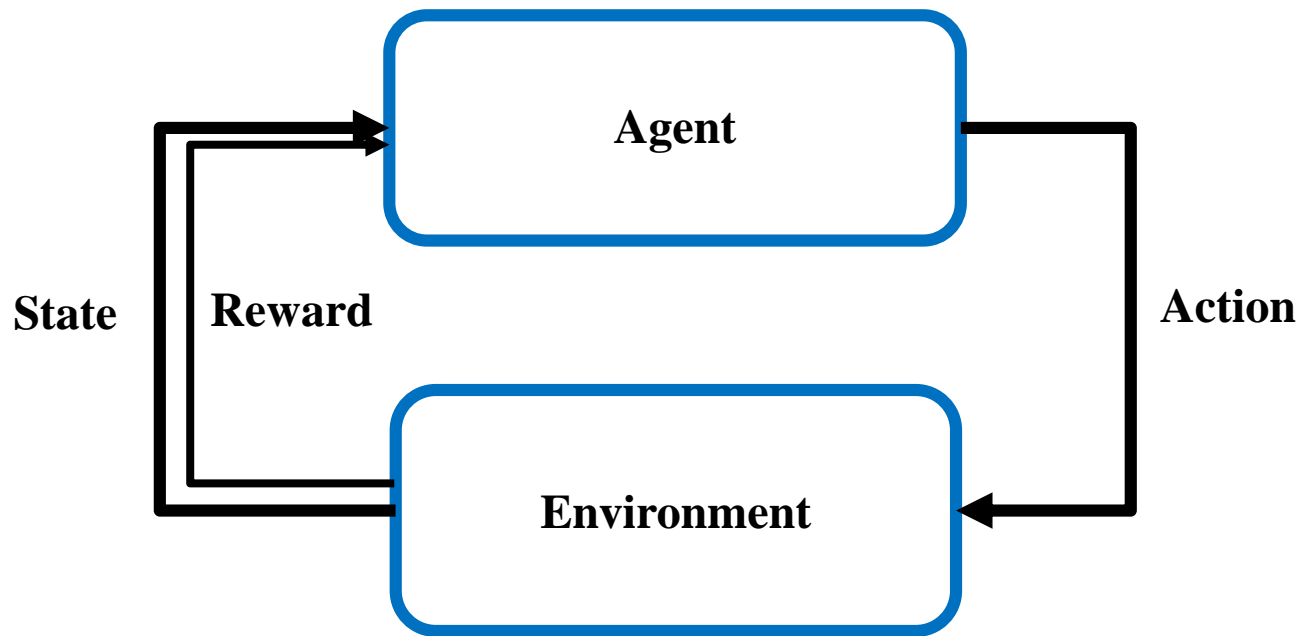
Reinforcement Learning

- A **computational approach** to **learning from interaction**
 - Explore designs for machines that are effective in
 - ▶ solving learning problems of scientific or economic interest,
 - ▶ evaluating the designs through mathematical analysis or computational experiments.
 - Focus on **goal-directed learning from interaction**, when compared with other approaches to machine learning.
 - The learner must discover which actions yield the most reward by trying them.
 - ▶ Two characteristics: most important distinguishing features of reinforcement learning.
 - **trial-and-error search**
 - **delayed reward**



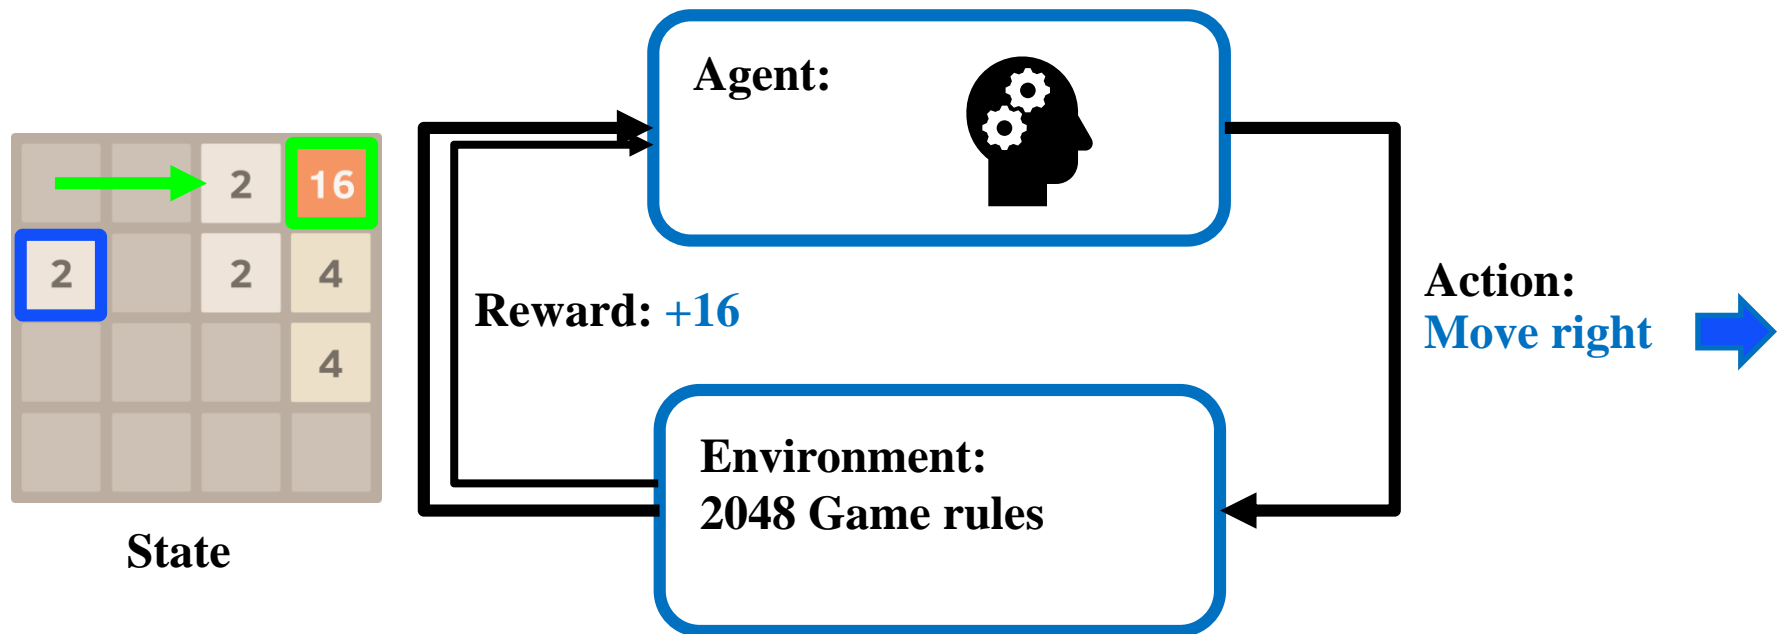
Agent-Environment Interaction Framework

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)



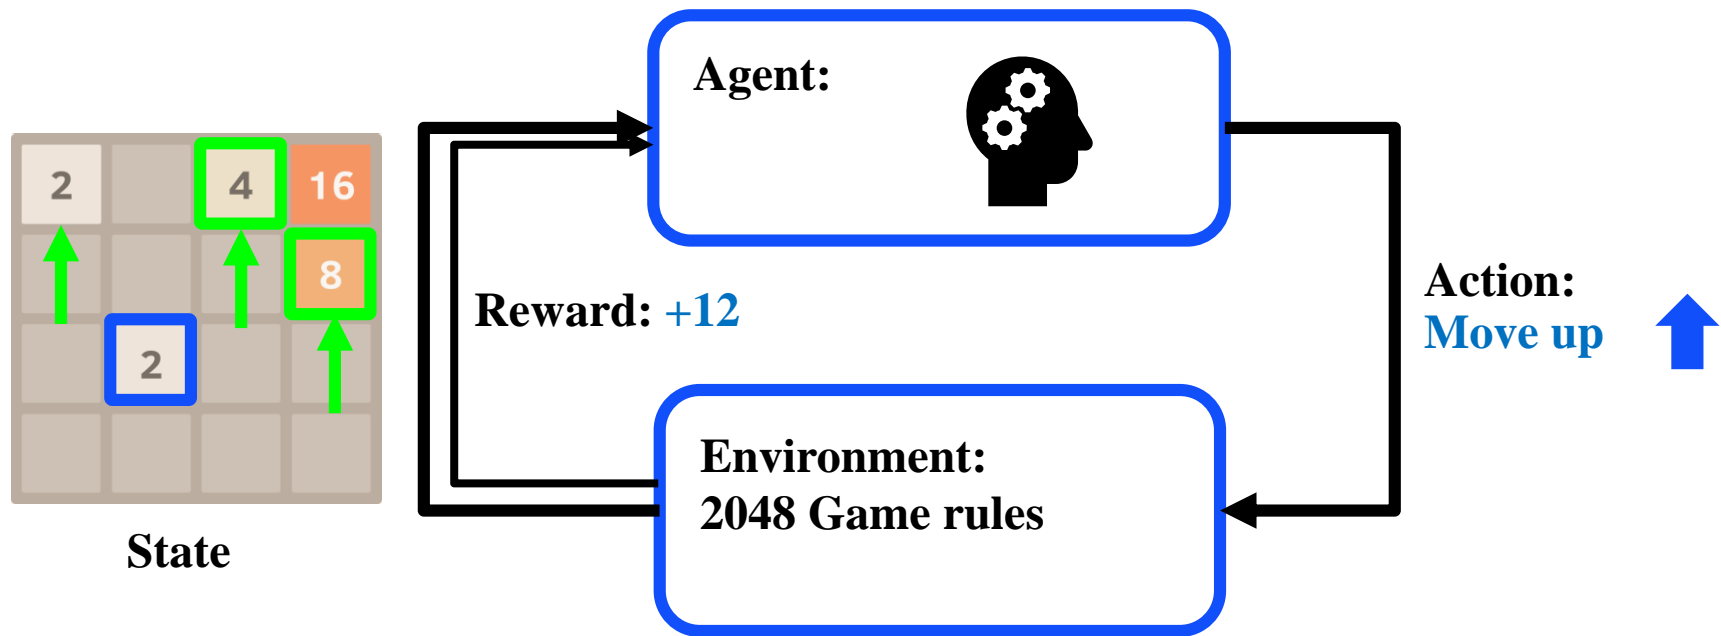
Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)



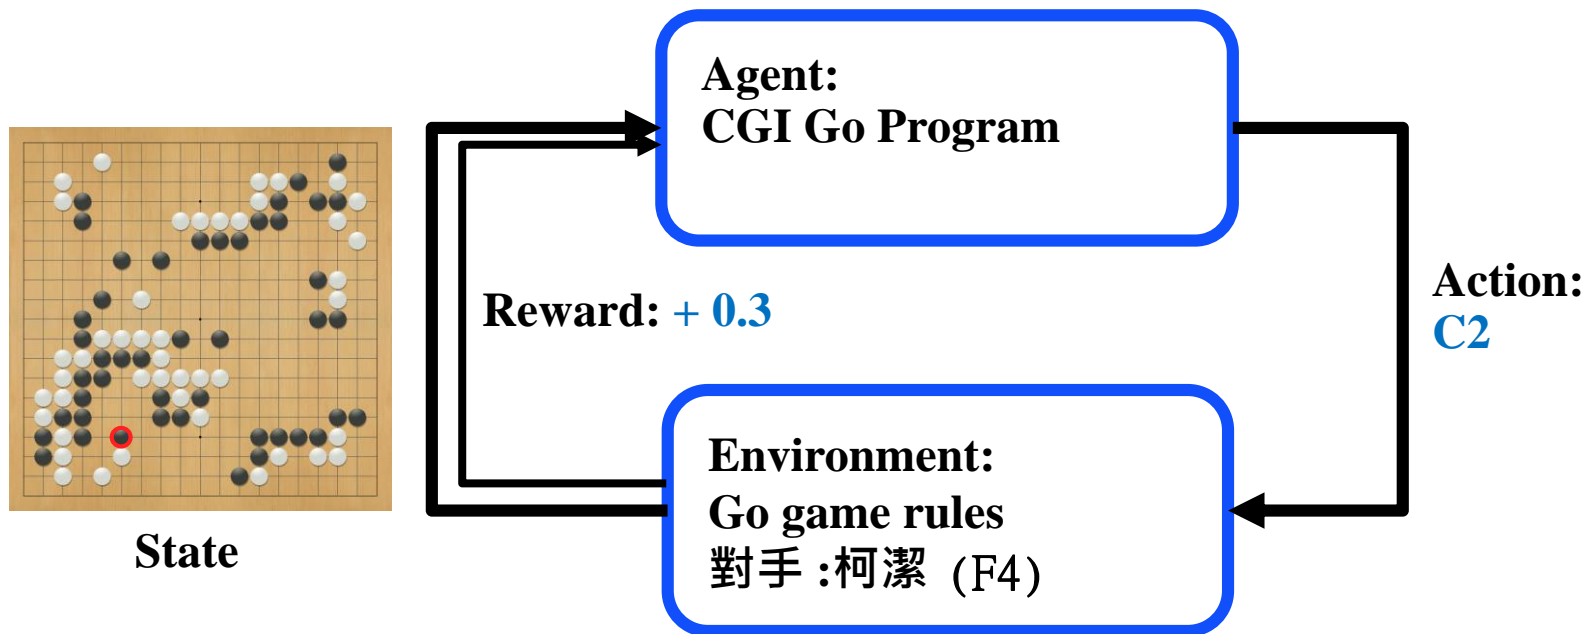
Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)




Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)




Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)



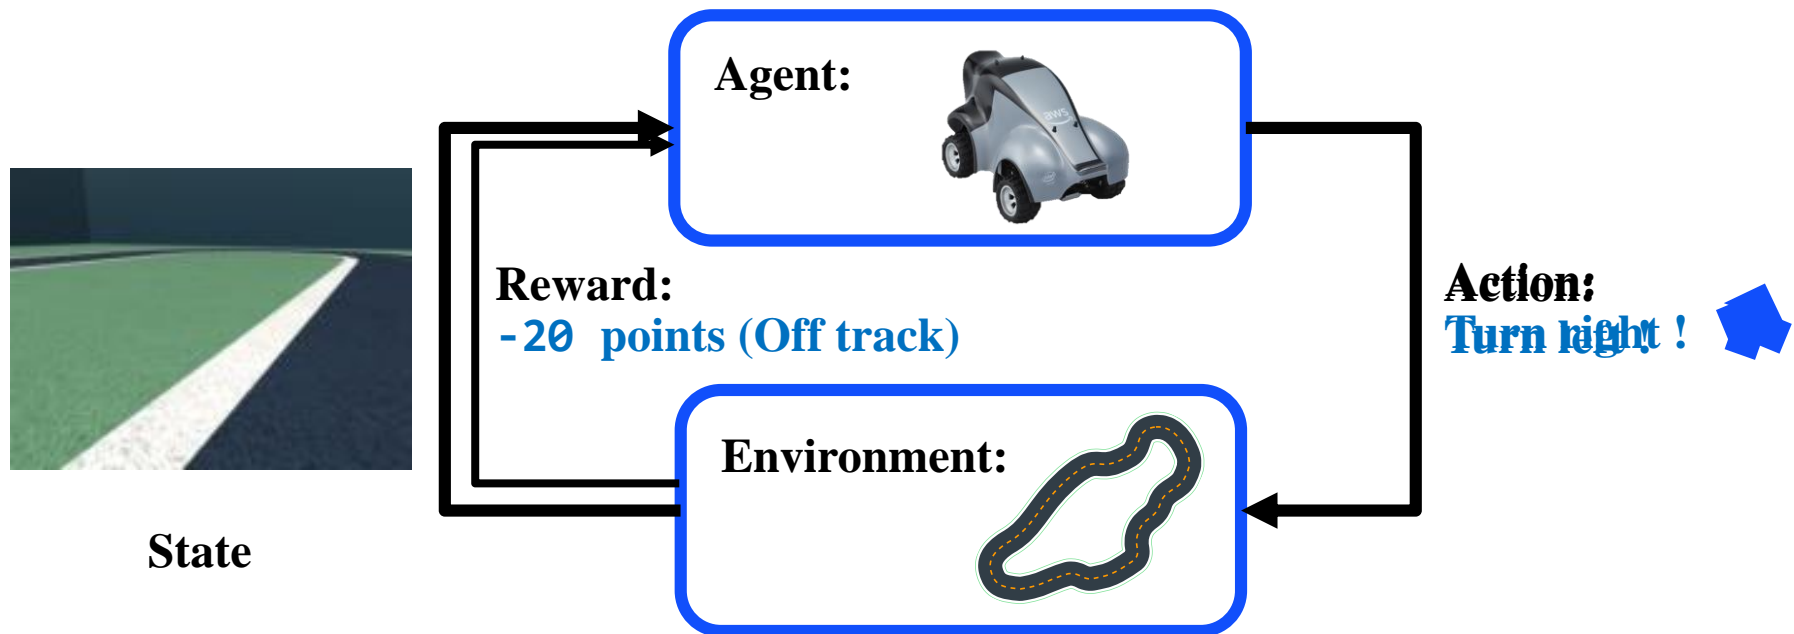
Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)



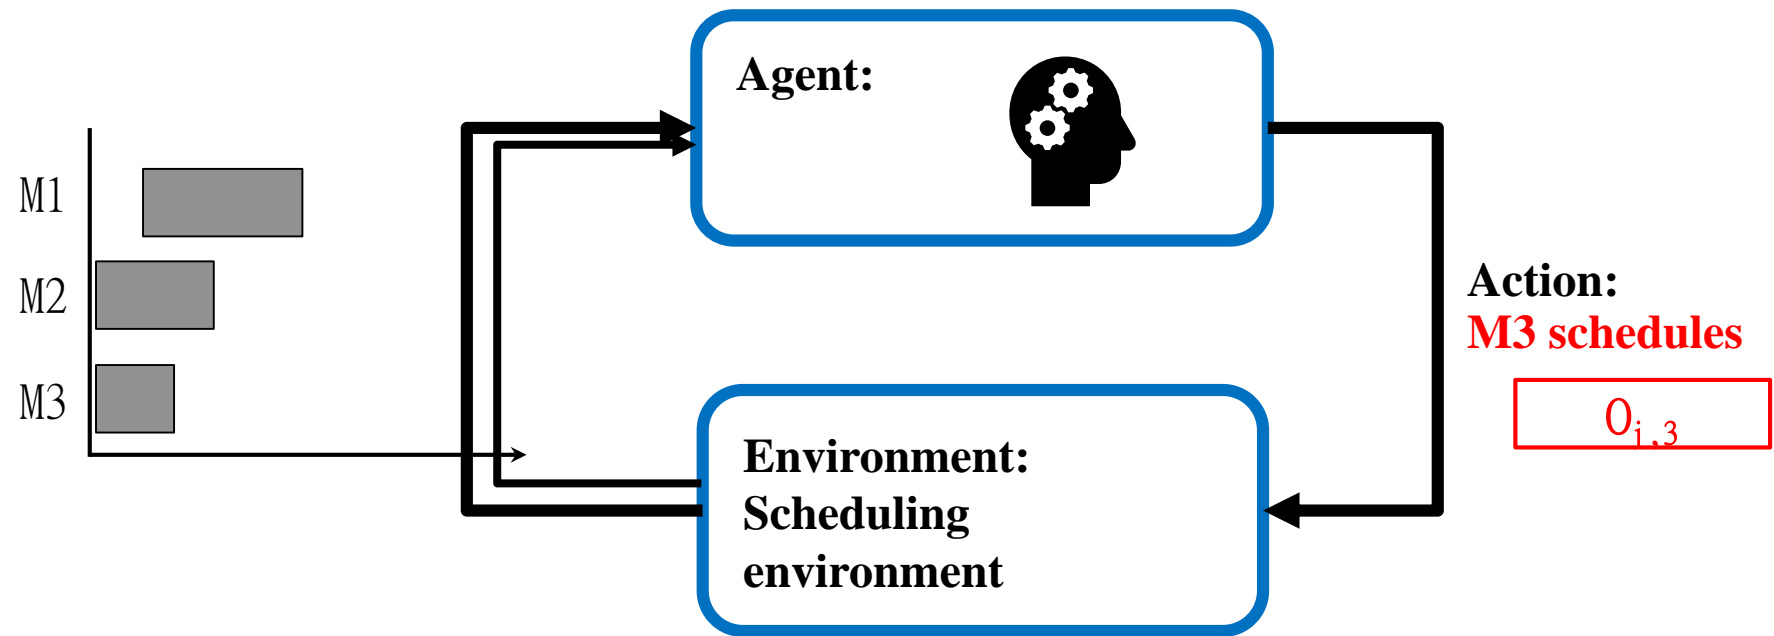
Reinforcement Learning (RL)

- A kind of AI **computational approach** to **learning from interaction**
- Agent-Environment Interaction Framework (代理者-環境 互動框架)



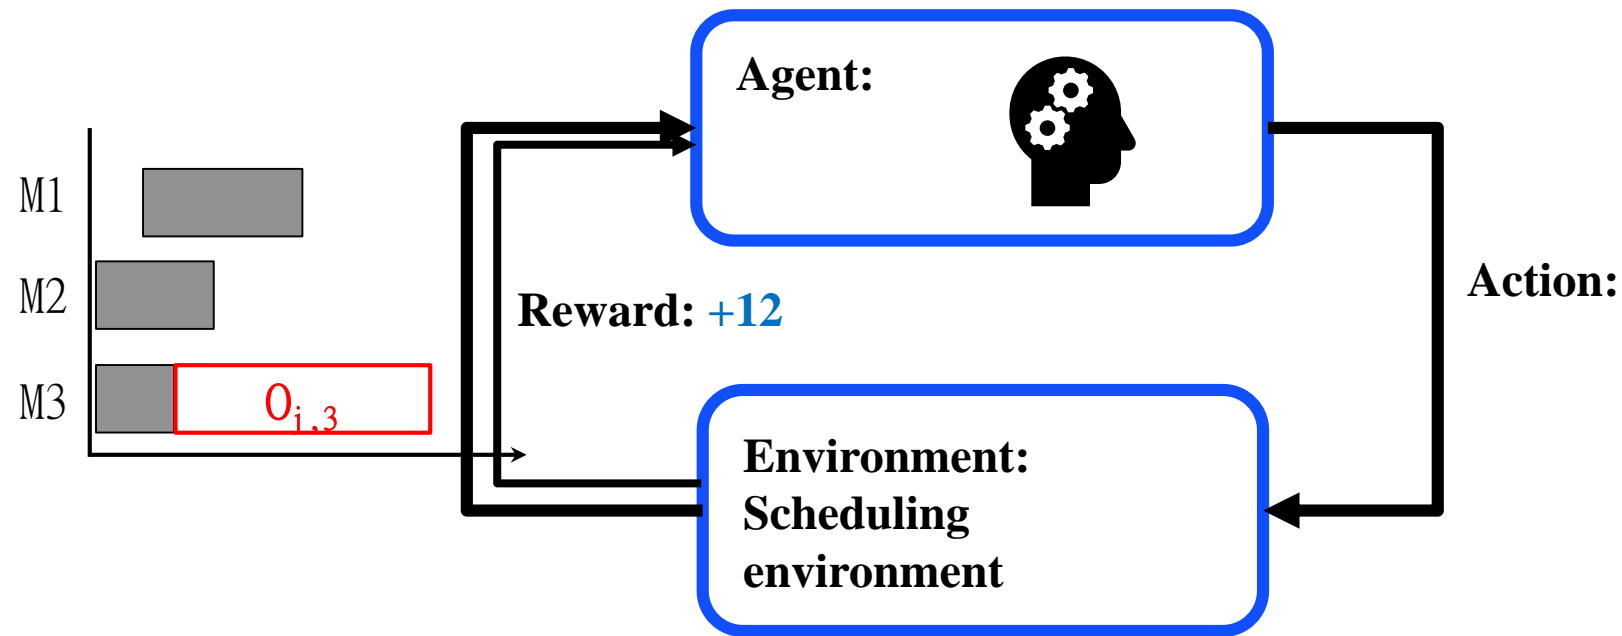
Reinforcement Learning (RL)

- A kind of AI **computational approach** to learn **from interaction**
- Agent-Environment Interaction Framework

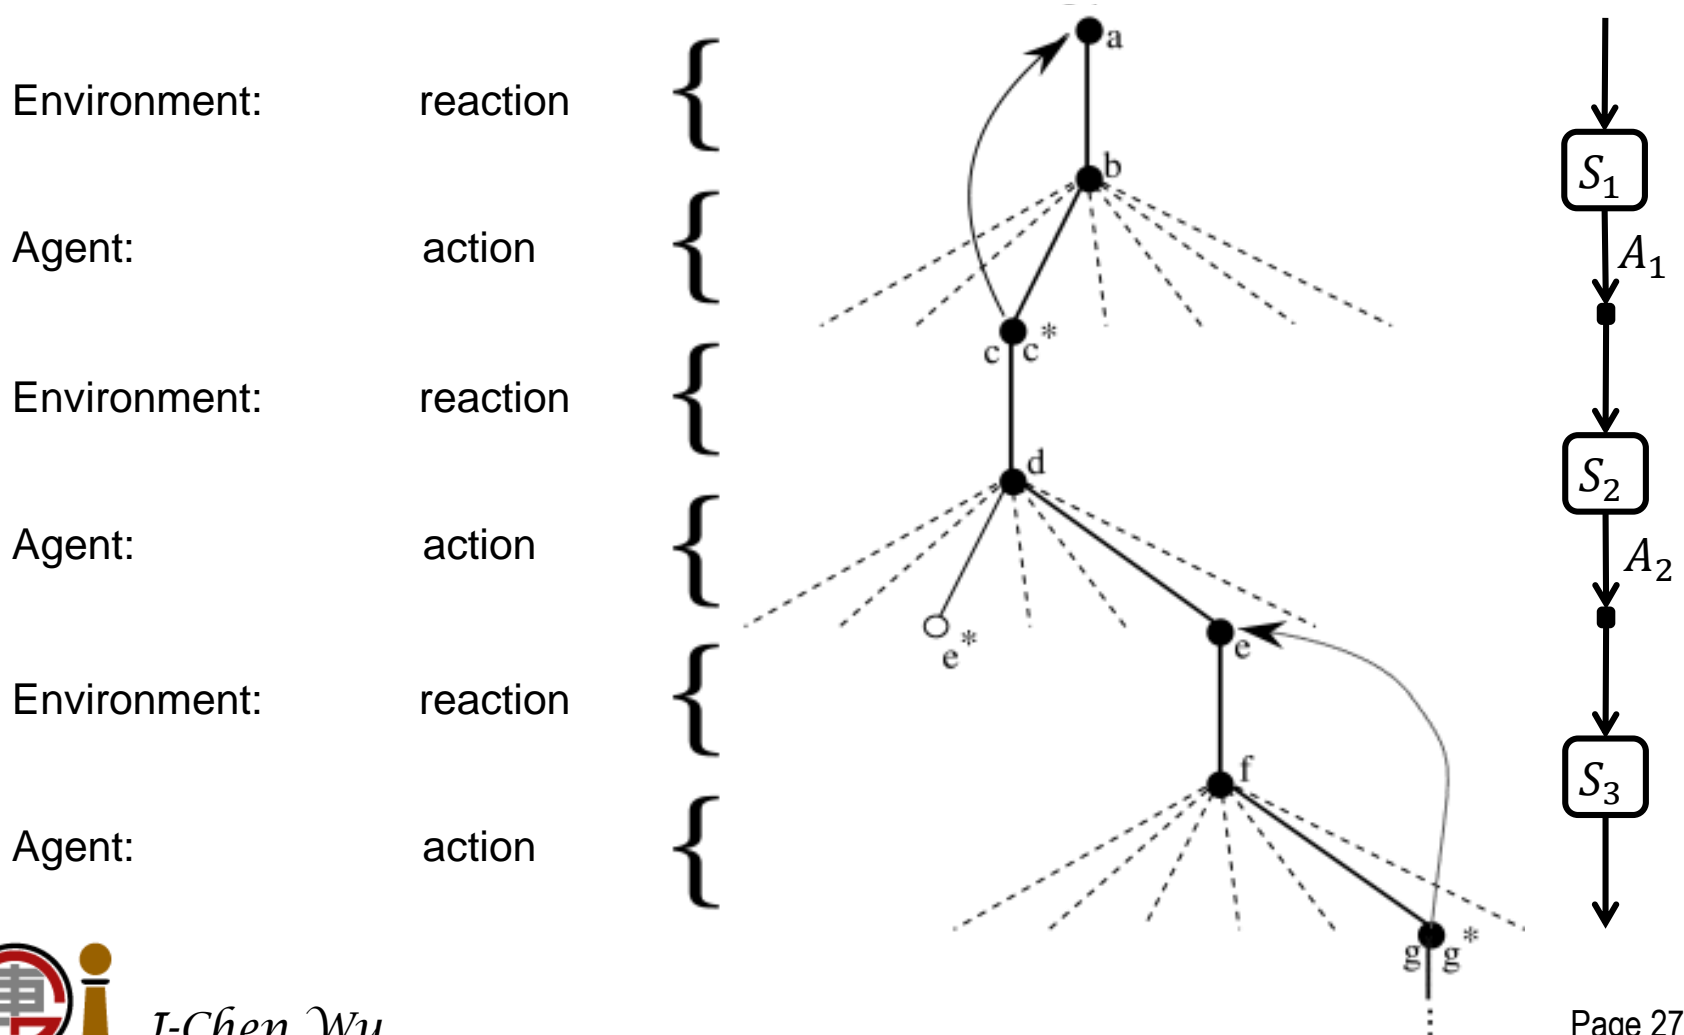


Reinforcement Learning (RL)

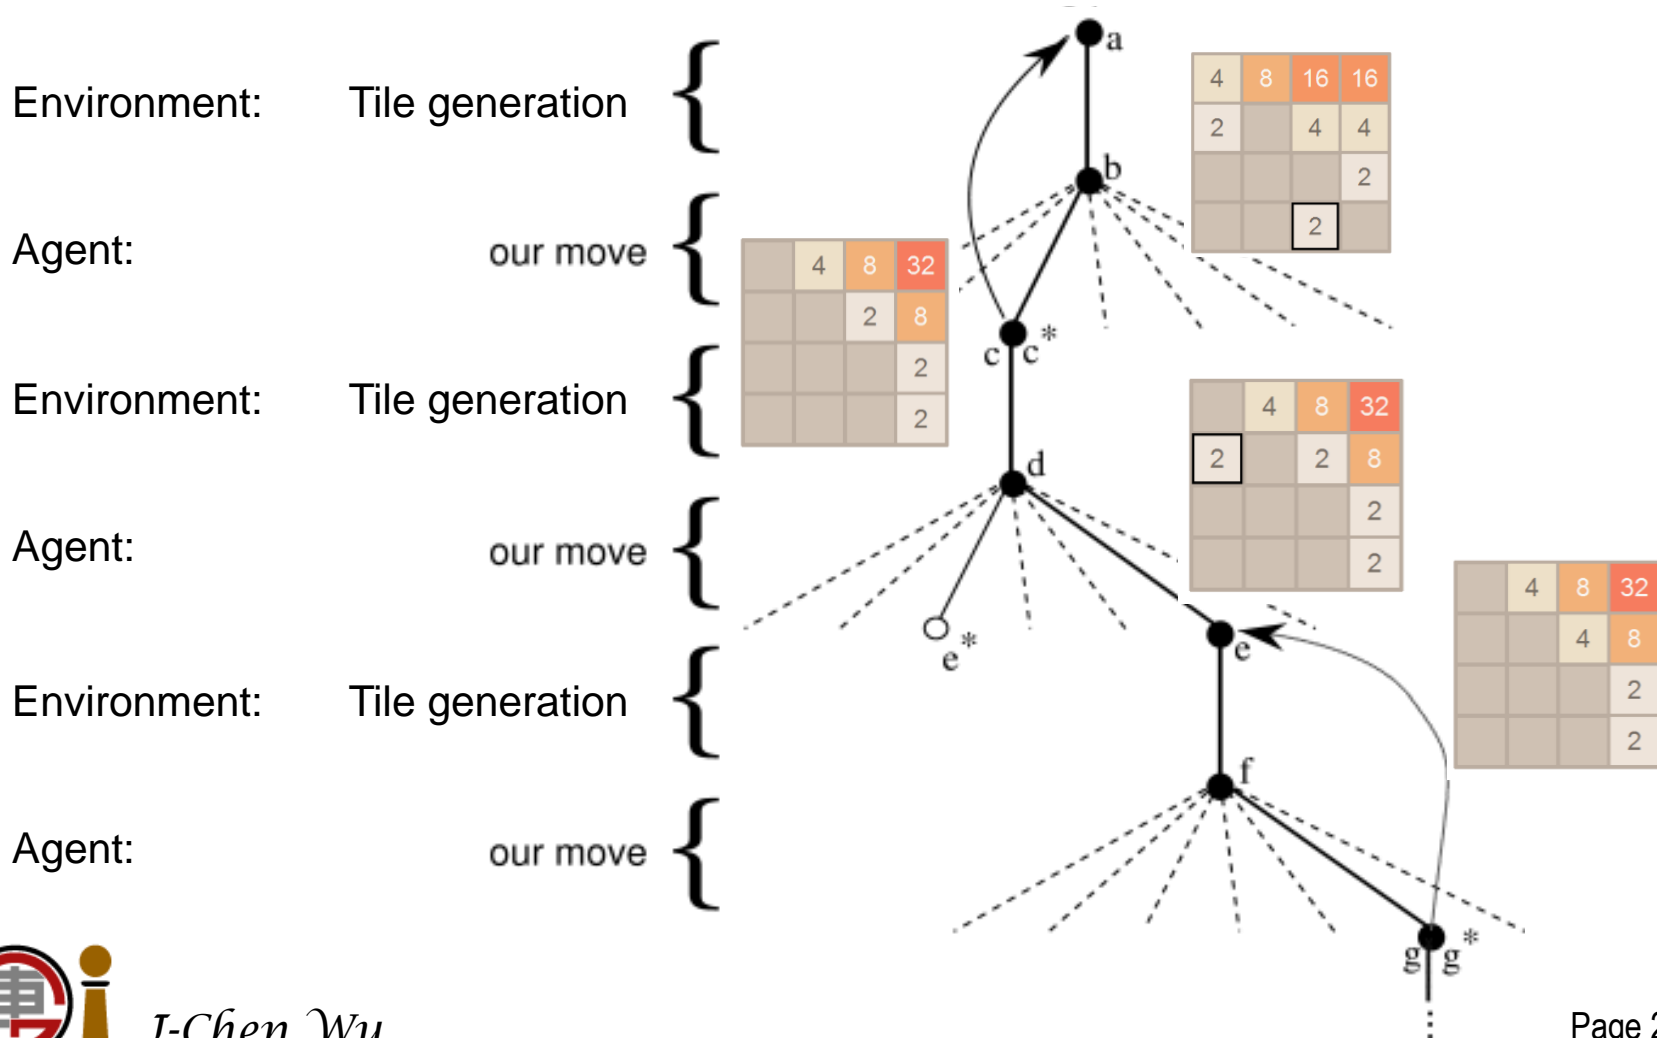
- A kind of AI **computational approach** to learn **from interaction**
- Agent-Environment Interaction Framework



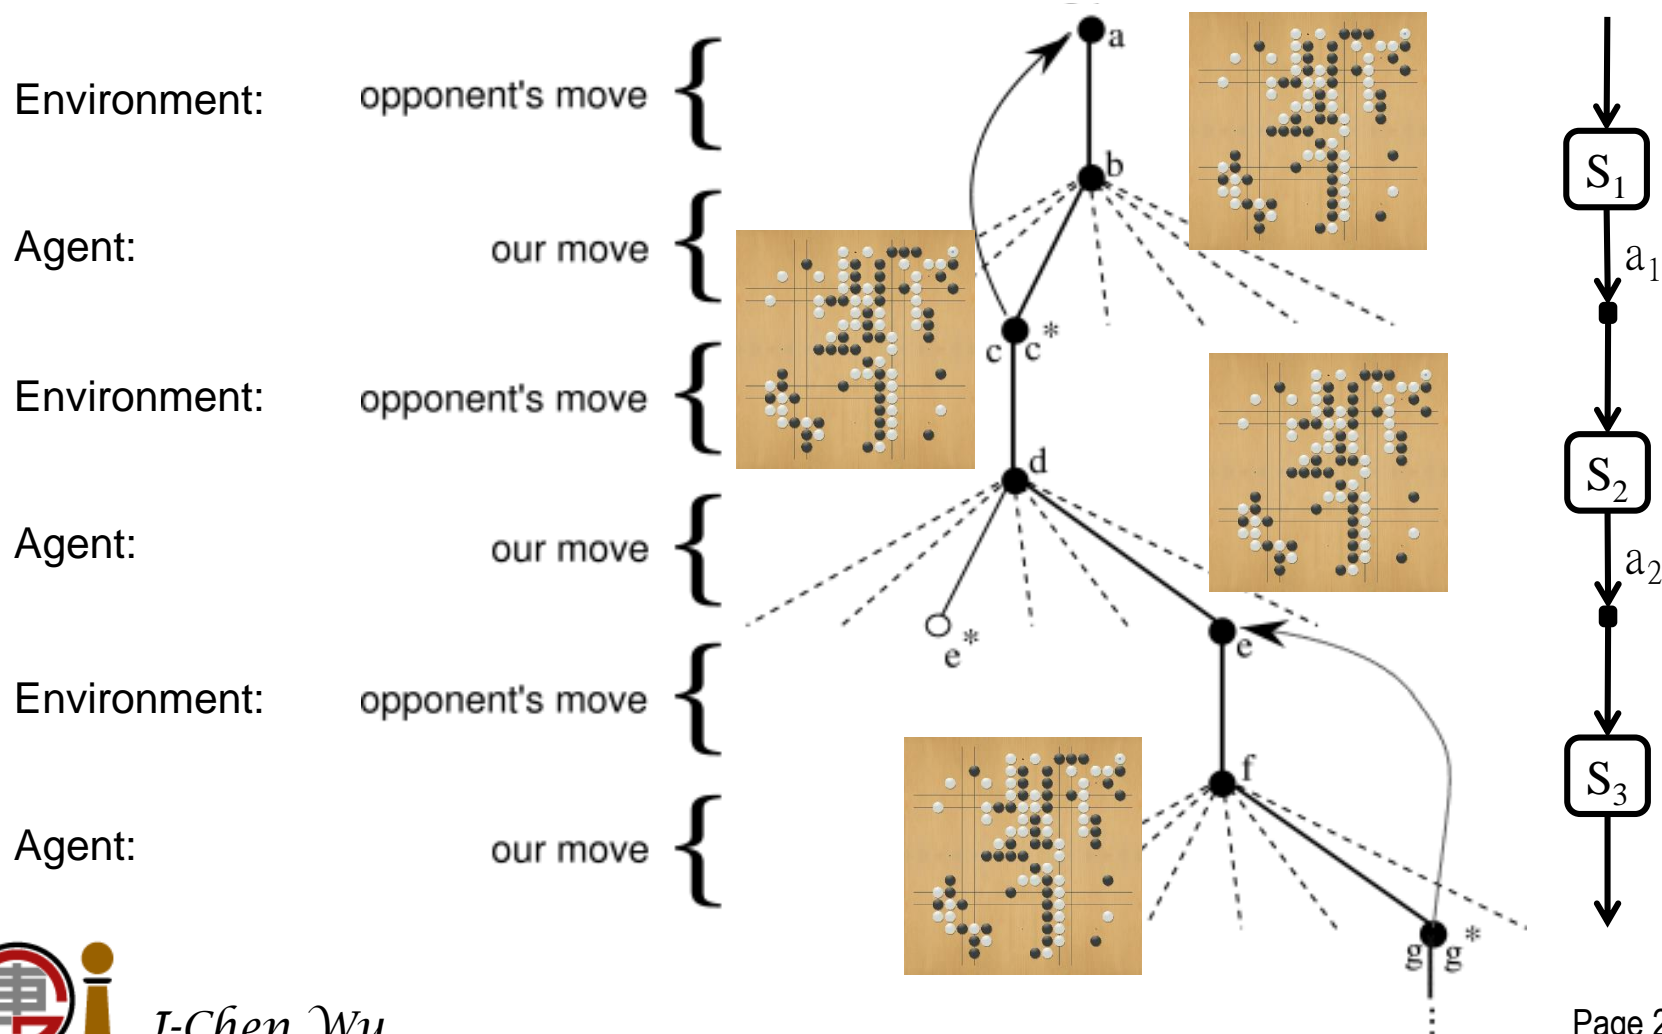
States and Actions in the Framework



2048



Go



Robot

Environment: Dynamics


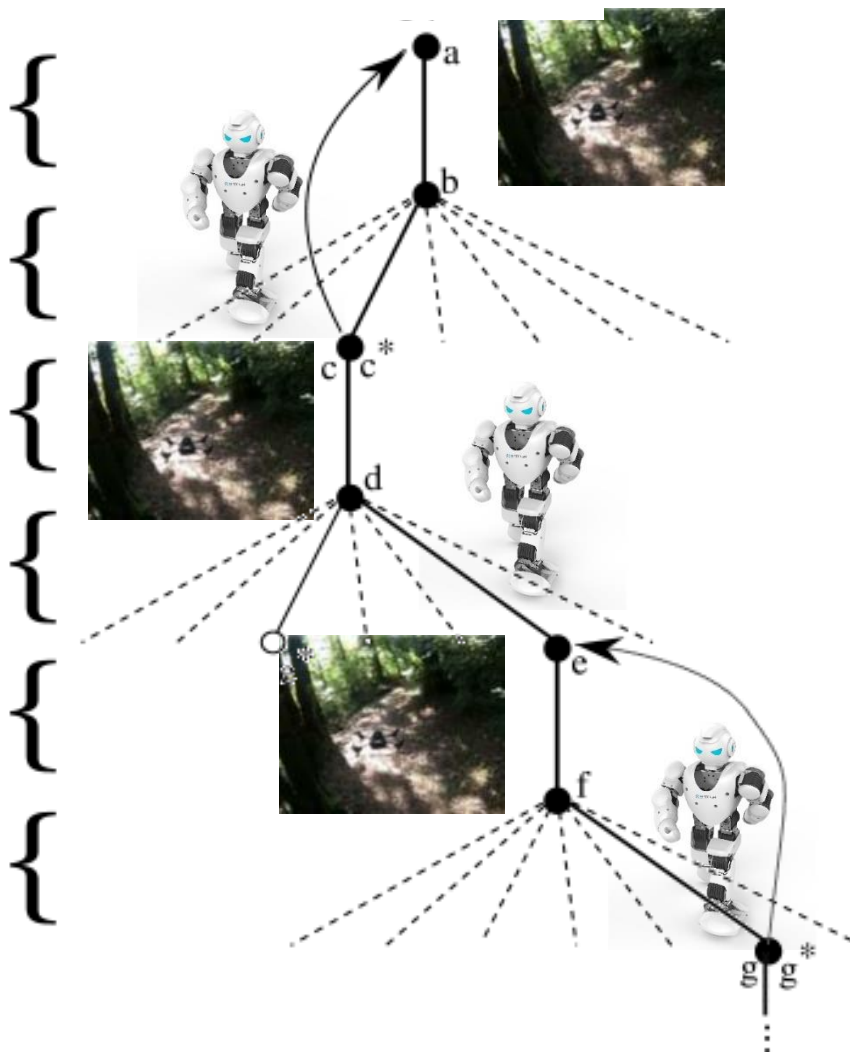
Agent: Navigate

Environment: Dynamics


Agent: Navigate

Environment: Dynamics



Agent: Navigate




More Example: Flappy Bird



More Toy Example: Pendulum




More Example: RL Demo (DDPG)




Markov Decision Processes (MDP)

- A (Finite) **Markov Decision Process** is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
 - \mathcal{S} is a (finite) set of states
 - \mathcal{A} is a (finite) set of actions
 - \mathcal{P} is a state transition probability matrix (part of the environment),
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$
 - \mathcal{R} is a reward function,
 $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$
 - γ is a discount factor $\gamma \in [0, 1]$.



Markov Property



- An **episode**: (assuming finite and MDP here for simplicity)


- States: S_i
 - ▶ Initial state: S_1
 - ▶ Current state: S_t
 - ▶ End state: S_T (not necessarily required)
- Actions: A_i
- **History**: $H_t = (S_1, A_1, R_2, S_2, A_2, R_3, S_3, \dots, R_t)$

- Markov Property:

- “The future is independent of the past given the present”
- A state S_t is **Markov** if and only if
$$\mathbb{P}[S_{t+1} | S_t] = \mathbb{P}[S_{t+1} | S_1, \dots, S_t]$$




Episode and Space



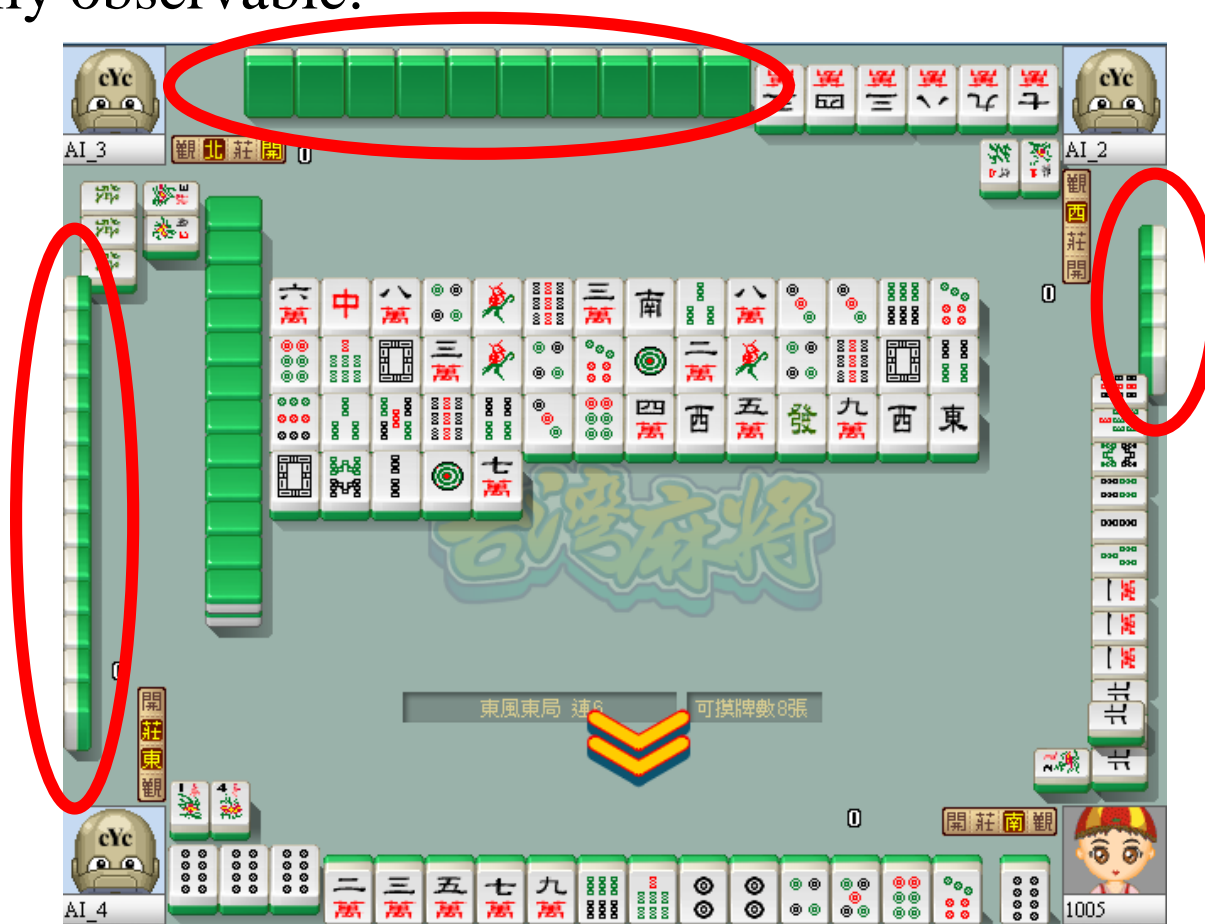
Environment State vs. Agent State

- The **environment state S_t^e** :
 - the environment's private representation
 - ▶ i.e. whatever data the environment uses to pick the next observation/reward
 - The environment state is not necessarily visible to the agent
 - ▶ Even if S_t^e is visible, it may contain irrelevant information
- The **agent state S_t^a** :
 - The agent's internal representation
 - ▶ i.e. whatever information the agent uses to pick the next action
 - ▶ i.e. it is the information used by reinforcement learning algorithms
 - It can be any function of history:
$$S_t^a = f(H_t)$$
- **Partially Observable**: (not discussed here)
 - When $S_t^a \neq S_t^e$




Example: Mahjong


- Partially observable:



Rewards



- A reward R_t is a **scalar feedback** signal
 - Indicates how well agent is doing at step t
 - The agent's job is to maximize cumulative reward
 - Reinforcement learning is based on the **reward hypothesis**
 - Example: (2048)




Definition (Reward Hypothesis)

- All goals can be described by the maximization of expected cumulative reward

Sequential Decision Making

- Goal:
 - Select actions to maximize total future reward
- Maximize $R_{t+1} + R_{t+2} + \dots + R_T$
 - assuming time = t .



- Notes:
 - Actions may have long term consequences
 - Reward may be delayed
 - It may be better to sacrifice immediate reward to gain more long-term reward



Sequential Decision Making – Examples

- Examples:

- In 2048, establish a sequence of $(2^t, 2^{t-1}, 2^{t-2}, \dots)$
- In chess, block opponent moves to help winning chances many moves from now.
- In a financial investment, may take months to mature
- In robotics, refuel a helicopter to prevent a crash.

2	32768	8192	4096
16384	1024	512	256
2048	32	64	128
16	16	2	4

Return

Definition

- The return G_t is the total discounted reward from time-step t .

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Notes:

- The discount $\gamma \in [0, 1]$ is the present value of future rewards
- The value of receiving reward R is diminishing
 - $\gamma^k R$, after $k + 1$ time-steps.
- This values immediate reward above delayed reward.
- Discount:
 - γ close to 0 leads to "myopic" evaluation
 - γ close to 1 leads to "far-sighted" evaluation
 - Important for infinite episodes.




Major Components of an RL Agent


- **Value function**: how good is each state and/or action
- **Policy**: agent's behavior function
- **Model**: agent's representation of the environment

Policy

- A policy is the agent's behavior
 - It is a map from state to action,
- Policy types:
 - Deterministic policy: $a = \pi(s_i)$
 - Stochastic policy: $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$
 - ▶ Sometimes, written in $\pi(s, a)$.
- Examples:
 - In 2048: Up/down/left/right
 - In robotics: angle/force/...




Value Function



- A value function is
a prediction of future reward
 - Used to evaluate the goodness/badness of states
 - ▶ therefore to select between actions.
 - Return $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- Types of value functions under policy π :
 - **State value function**: the expected return from s .
$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$
$$= \mathbb{E}_\pi[G_t \mid S_t = s]$$
 - **Q-Value function**: the expected return from s taking action a .
$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a]$$
- Examples:
 - In 2048, the expected score from a board S_t .

Model



- A **model** predicts

what the environment will do next

- \mathcal{P} is a state transition probability matrix,

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$
 - ▶ predicts the next state
- \mathcal{R} is a reward function,

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$
 - ▶ predicts the next (immediate) reward

- Examples:

- In 2048:
 - ▶ After a move, \mathcal{P} is to generate a tile randomly as follows:
 - 2-tile: with probability of 9/10
 - 4-tile: with probability of 1/10



Categorizing RL Agents (Policy & Value)

- Value Based
 - No Policy (Implicit)
 - Value Function
- Policy Based
 - Policy
 - No Value Function (Implicit)
- Actor Critic
 - Policy
 - Value Function



